

CS 315 - Day 34, Pattern Matching

The broad topic we begin discussing today are String algorithms.

In particular, we will discuss:

- String and Pattern Matching algorithms
- Text **Similarity** testing (used in DNA sequencing)
- File compression
- Data structures and algorithms used in search engines

Document processing is one of the dominant functions of technology.

We start with pattern matching, which formally defined is:

Given a **pattern** and a **text**, we want to find all occurrences of the pattern in the text. Many applications, including searching for text in a document, finding web pages relevant to queries, and searching for patterns in DNA sequences.

Formally:

Input: The text is an array $T[1:n]$, and the pattern is an array $P[1:m]$, where $m \leq n$. The elements of P and T are drawn from a finite **alphabet** Σ . Examples: $\Sigma = \{0, 1\}$, Σ is the ASCII characters, or $\Sigma = \{A, C, G, T\}$ for DNA matching. Call the elements of P and T **characters**.

Output: All amounts that we have to shift P to match it with characters of T . Say that P **occurs with shifts in T** if $0 \leq s \leq n-m$ and $T[s+1:s+m] = P[1:m]$. We want to find all valid shifts.

This is the naïve/brute-force approach.

Σ is the alphabet of all the possible characters, T is the input string to search, and P is the pattern searched for.

Example: $\Sigma = \{A, C, G, T\}$, $T = \text{GTAACAGTAAACG}$, $P = \text{AAC}$.

Just try each shift.

NAIVE-STRING-MATCHER(T, P, n, m)

```
for s = 0 to n - m
    if P[1:m] == T[s + 1:s + m]
        print "Pattern occurs with shift" s
```

The shift s is the number of characters to “skip over,” so P starts at $s + 1$ in T . As defined above, P is m characters long and T is n characters long.

In the example above, we have shifts $s = 2$ and $s = 9$.


G	T	A	A	C	A	G	T	A	A	A	C	G
1	2	3	4	5	6	7	8	9	10	11	12	13

Time: Tries $n - m + 1$ shift amounts, each taking $O(m)$ time, so $O((n - m + 1)m)$.

This bound is tight in the worst case, such as when the text is n As and the pattern is m As. No preprocessing needed.

This algorithm is not efficient. It throws away valuable information.

Brute Force?



It gets expensive if you “almost” find the string, over and over again.

information learned during the search. The first algorithm makes use of a structure we explored in MA 116 – Discrete Structures, a Finite-State-Automaton (FSA).


Formally: A finite automaton M is a 5-tuple $(Q, q_0, A, \Sigma, \delta)$ where

- Q is a finite set of *states*,
- $q_0 \in Q$ is the *start state*,
- $A \subseteq Q$ is a distinguished set of *accepting states*,
- Σ is a finite *input alphabet*,
- δ is a function $Q \times \Sigma \rightarrow Q$, called the *transition function* of M .

The FA begins in state q_0 and reads the characters of the input string one at a time. When in state q and reading character a , it moves to state $\delta(q, a)$. If the current state is in A , then the FA has *accepted* the string read so far. An input that is not accepted is *rejected*.

For string matching, the FA has $m + 1$ states, numbered 0 to m . The FA starts in state 0. When it's in state k , the k most recent text characters read match the first k characters of the pattern. When the FA gets to state m , it has found a match.

Example: Use the alphabet $\Sigma = \{A, C, G, T\}$ for DNA sequencing. If $P = ACACAGA$ with $m = 7$, then the FA is [leave this figure on the board]



The horizontal spine has edges labeled with P . Whenever P occurs in the text, the FA moves right, to the next state along the spine. When it reaches the rightmost state, it has found a match.

The transition function δ is defined for all states $q \in Q$ and all characters $a \in \Sigma$. Missing arrows are assumed to be transitions to state 0.

When a character from the text does not make progress toward a match, the FA moves to a lower-numbered state or stays in the same state. The only arrows pointing to the right are along the spine.

We'll see how to compute the transition function δ later. Assuming that we have it, here's how to perform string matching.

FA-MATCHER(T, δ, n, m)

```

 $q = 0$ 
for  $i = 1$  to  $n$ 
   $q = \delta(q, T[i])$ 
  if  $q == m$ 
    print "Pattern occurs with shift"  $i - m$ 
```


So, can you describe the **worst-case** for the Naïve-String-Matcher algorithm?

Almost finding the string each loop, and the length of the Pattern P is $m = \left\lfloor \frac{n}{2} \right\rfloor$

So, we will explore two algorithms which are better because they pre-process or save

In-class w/assigned groups of four (4), write 0. and 1. on board as ground rule(s):

0. No devices.
1. Intro: Name, POE, Icebreaker.
2. Pick someone on the group who will report out.
3. Create the State-Transition-Diagram for $\Sigma = \{a, b\}$, $P = aabab$
- Hint: $m + 1$ states
4. Stretch goal, create the Next-State-Table.



So, we could just create a Next-State-Table directly from the State-Transition-Diagram.

It would have $m + 1$ rows and $|\Sigma|$ columns, and is called $\delta(\text{current-state}, \text{input-char})$, and outputs the next state:

State Transition Table

State	a	b	P
0	$\delta(0,a) = \sigma(P_0a) = \sigma(a) = 1$	$\delta(0,b) = \sigma(P_0b) = \sigma(b) = 0$	a
1	$\delta(1,a) = \sigma(P_1a) = \sigma(aa) = 2$	$\delta(1,b) = \sigma(P_1b) = \sigma(ab) = 0$	a
2	$\delta(2,a) = \sigma(P_2a) = \sigma(aaa) = 2$	$\delta(2,b) = \sigma(P_2b) = \sigma(aab) = 3$	b
3	$\delta(3,a) = \sigma(P_3a) = \sigma(aaba) = 4$	$\delta(3,b) = \sigma(P_3b) = \sigma(aabb) = 0$	a
4	$\delta(4,a) = \sigma(P_4a) = \sigma(aabaa) = 2$	$\delta(4,b) = \sigma(P_4b) = \sigma(aabab) = 5$	b
5	$\delta(5,a) = \sigma(P_5a) = \sigma(aababa) = 1$	$\delta(5,b) = \sigma(P_5b) = \sigma(aababb) = 0$	

$\sigma(P_i a)$ = "length of the longest prefix of P which is a suffix of $P_i a$ "

But we should be able to automate the creation of the table:

How do we build the finite automaton? We already know that it has states 0 through m , the start state is 0, and the only accepting state is m . The hard part is the transition function δ . The idea:

When the FA is in state k , the k most recent characters read from the text are the first k characters in the pattern.

A *prefix* $P[:i]$ of a string P is a substring consisting of the first i characters of P . A *suffix* is a substring consisting of characters from the end. For a string X and a character a , denote concatenating a onto X by Xa .

So, in state k , we've most recently read $P[:k]$ in the text. We look at the next character a of the text, so now we've read $P[:k]a$. How long a prefix of P have we just read?

Find the longest prefix of P that is also a suffix of $P[:k]a$. Then $\delta(k, a)$ should be the length of this longest prefix.

In our example, $\delta(q, \text{input} - \text{char}) = \sigma(P[:q] \mid \text{input} - \text{char})$

State, $q = 2$, $\text{input} - \text{char} = a \rightarrow \delta(2, a) = P[:2] \mid a = aa \mid a = aaa$
 $q = 2$, $\text{input} - \text{char} = b \rightarrow \delta(2, b) = P[:2] \mid b = aa \mid b = aab$

COMPUTE-TRANSITION-FUNCTION(P, Σ, m)

```
1  for q = 0 to m
2      for each character a ∈ Σ
3          k = min {m, q + 1}
4          while P[:k] is not a suffix of P[:q]a
5              k = k - 1
6          δ(q, a) = k
7  return δ
```

Preprocessing time:

- The outermost **for** loop iterates $m + 1$ times.
- The middle loop iterates $|\Sigma|$ times.
- The innermost **while** loop iterates at most $m + 1$ times.
- Each suffix check in the **while** loop test examines at most m characters of the pattern (since $k \leq m$). Therefore, each suffix check takes $O(m)$ time.
- Total preprocessing time: $O(m^3 |\Sigma|)$.

Therefore, with a finite automaton, we can perform string matching with preprocessing time $O(m^3 |\Sigma|)$ and matching time $\Theta(n)$. It's possible to get the preprocessing time down to $O(m |\Sigma|)$.

The downside is the expense to pre-process and create the State Transition Table.

(# symbols in alphabet) * (length of P)³

But we can do even better, reducing the pre-processing time to $\theta(m)$ and keeping the matching time at $\theta(n)$. It is the Knuth-Morris-Pratt (KMP) algorithm, and it uses an auxillary array $\pi[1..m]$, which allows re-use of comparisons.

This array encapsulates knowledge about how the pattern matches against shifts of itself. We avoid (1) testing useless shifts and (2) pre-computing the Next-State-Table δ .

T back ababababcb ab
 $\xrightarrow{s \rightarrow}$ aba ba ca P
First q=5 characters match

Using only our knowledge of the 3 mismatched characters, we know a shift of $s+1$ would be invalid. A shift of $s+2$ might work

b a c b a b a b a a b c b a b
 $\xrightarrow{s+2}$ a b a b a c a
(k=3)

In the example above, the longest prefix of P that is also a proper suffix of P_5 is P_3

ab a ba $\xrightarrow{k=3}$
 a ba P_3)

$\{\alpha\} \subseteq \{\alpha\}$
 $\{\alpha\} \not\subseteq \{\alpha\}$
 \subseteq vs \subset

$$\pi[5] = 3$$

Prefix for
 P_3

$$\pi[g] = \max \{ k : k < g \text{ and } P_k \sqsupseteq P_g \}$$

$\pi[g]$ is the length of the longest

prefix of P that is a proper suffix
 of P_g .

$(k < g)$ $(k \leq g)$


If g characters have successfully matched at shift s , the next potentially valid shift is @

$$s + (g - \pi[g])$$

ex:

i	1	2	3	4	5	6	7	8	9	10
$P[i]$	a	b	a	b	a	b	c	a		
$\pi[i]$	0	0	1	2	3	4	5	6	0	1

	1	2	3	4	5	6
P[i]	a	b	a	c	a	b
T[i]	0	0	1	0	1	2



Compute Prefix Function

```
Compute-Prefix-Function (P)
m = P.length
Pi[1] = 0
k = 0
for q=2 to m
    while k>0 and P[k+1] != P[q]
        k = Pi[k]
    if P[k+1] = P[q]
        k++
    Pi[q] = k
return Pi
```

KMP Matcher

```
KMP-Matcher (T, P)
n = T.length
m = P.length
Pi = Compute-Prefix-Function
q = 0
for i=1 to n
    while q>0 and P[q+1] != T[i]
        q = Pi[q]
    if P[q+1] = T[i]
        q++
    if q = m
        print "Pattern occurs w/shift" i-m
        q = Pi[q] // look for the next match
```